PRESS RELEASE

Kiro Power: AI-Powered AWS Observability Plugin for Claude Code

Launch Date: April 29, 2026

Target Customer: SREs, DevOps engineers, and platform teams running production services on AWS Application Signals

SEATTLE, WA -- April 29, 2026 -- Today, awslabs announced the general availability of Kiro Power, an open-source Claude Code plugin that turns natural language into structured AWS APM investigations. Kiro Power connects to four existing AWS MCP servers (CloudWatch, Application Signals, CloudTrail, and AWS Documentation) and encodes a repeatable, six-phase investigation method that produces cited, verifiable artifacts every time. SREs and service-owning developers can now go from page to root-cause hypothesis in under five minutes without leaving their IDE.

The Problem

When an alarm fires at 3 AM, on-call engineers face the same costly sequence: open the CloudWatch console, cross-reference metrics with traces, search logs for error patterns, check CloudTrail for recent deployments, and mentally correlate signals across four or five browser tabs. This context-switching is slow, error-prone, and depends on tribal knowledge about which metrics matter for which services. The result is inconsistent investigations where critical evidence is missed, incident channels receive incomplete updates, and the same regressions are re-investigated week after week because there is no persistent memory of prior root causes.

The Solution

Kiro Power eliminates this context-switching by encoding a fixed, six-phase investigation method on top of the four AWS MCP servers. Engineers describe the problem in natural language ("got paged for HighCheckoutErrorRate" or "checkout p99 jumped from 200ms to 800ms") and the plugin executes a structured workflow: classify the issue, identify top contributors, gather metric/log/trace evidence, correlate CloudTrail changes, rank hypotheses with cited evidence, and follow downstream dependencies. Every investigation produces the same artifact shape -- a Service Health Card, SLO Breach Explainer, or Trace Waterfall Summary -- with a metadata footer citing every MCP tool call, a confidence level, and deep links directly into the CloudWatch console.

Customer Quote

"Before Kiro Power, our on-call engineers spent 15-20 minutes per incident just gathering context from different AWS console tabs. Now they describe the alert in plain English and get a cited, ranked hypothesis in under five minutes. The incident memory feature alone saved us from re-investigating the same Monday-morning DynamoDB throttling issue three weeks in a row. We have reduced our mean time to first hypothesis by 70 percent."

-- Jordan Patel, VP of Engineering, Meridian Commerce (fictional)

Getting Started

Kiro Power is free, open source (MIT license), and installs in under two minutes. Engineers add the plugin marketplace, run the install command, and configure their AWS profile and region. The plugin requires uv/uvx for MCP server management and valid AWS credentials with read-only IAM permissions for CloudWatch, Application Signals, X-Ray, and CloudTrail. No additional infrastructure, databases, or backing services are required. To start investigating, engineers type a natural language description of their issue or use one of twelve slash commands such as /cw-health-check, /cw-investigate-slo, or /cw-alarm-response.

Company Quote

"The MCP protocol gives AI agents access to tools. But tools alone are not enough -- agents need methods. Kiro Power demonstrates what becomes possible when you layer a domain-specific investigation method on top of standardized tool access: consistent, cited, verifiable results that an on-call engineer can trust at 3 AM. We are excited to see the community build on this pattern."

-- Dr. Amara Chen, Director of Developer Experience, AWS (fictional)

EXTERNAL FAQs

Product & Capabilities

Q: What is Kiro Power?

A: Kiro Power is an open-source Claude Code plugin that provides an AI-powered investigation method for AWS APM telemetry. It connects to four AWS MCP servers (CloudWatch, Application Signals, CloudTrail, and AWS Documentation) and encodes a repeatable six-phase workflow that produces structured, cited investigation artifacts.

Q: What can Kiro Power do?

A: Kiro Power supports three core investigation workflows (SLO breach, latency regression, error spike triage), two operational workflows (alarm response, SLO compliance reporting), and two production-readiness workflows (observability gap analysis, alerting design). It includes 20 skills, 12 slash commands, and 7 HTML artifact templates.

Q: Which AWS services does it work with?

A: Kiro Power works with AWS Application Signals, Amazon CloudWatch (metrics, alarms, Logs Insights), AWS X-Ray (distributed traces), and AWS CloudTrail (API audit trail). It is purpose-built for the AWS APM surface and does not cover billing, security-specific, or non-AWS APM platforms.

Q: What is the write-action safety model?

A: Kiro Power uses a 5-tier action safety model. Tier 1 (read-only telemetry) runs freely. Tier 2 (suggested actions) are described but not executed. Tier 3 (console deep-linked) directs users to the AWS console. Tier 4 (MCP-executable) requires explicit in-chat confirmation via a PreToolUse hook. Tier 5 (disallowed) actions such as deletions or IAM modifications are never executed.

Q: How does it ensure investigation quality?

A: Every investigation passes through a 6-check self-audit (investigation-validator) before output: metadata footer present, every claim cited to evidence, deep links valid, considered-and-ruled-out section present, burn-rate math correct, and confidence levels stated.

Getting Started

Q: How do I install Kiro Power?

A: Run two commands: add the plugin marketplace with /plugin marketplace add, then install with /plugin install aws-apm@aws-apm-plugins. The plugin also works in Cowork desktop via the plugin marketplace search. Total install time is under two minutes.

Q: What are the prerequisites?

A: You need uv/uvx installed (for MCP server management), configured AWS credentials (aws configure, AWS SSO, or environment variables), and Application Signals enabled in your AWS account. No additional infrastructure or databases are required.

Q: What IAM permissions are needed?

A: The recommended posture is read-only. The minimal IAM policy covers CloudWatch (GetMetricData, DescribeAlarms, StartQuery), Application Signals (ListServices, GetServiceLevelObjective, GetTraceSummaries), and CloudTrail (LookupEvents). Optional additive policies enable Logs Insights and CloudTrail advanced queries.

Q: Does it cost anything?

A: Kiro Power is free and open source under the MIT license. The only costs are standard AWS API call charges for CloudWatch, X-Ray, and CloudTrail queries, plus any Claude API usage from your existing Claude Code subscription.

Use Cases

Q: How does it help with incident response?

A: When an alarm fires, engineers type the alarm name or describe the issue in natural language. The plugin runs a structured six-phase investigation, correlates metrics with traces and CloudTrail changes, and produces a ranked hypothesis artifact with cited evidence and console deep links -- ready to paste into an incident channel.

Q: Can it help with proactive monitoring?

A: Yes. The /cw-health-check command produces a fleet-level dashboard across all Application Signals services. The /cw-obs-gaps command audits a codebase for logging, metrics, tracing, and error-handling gaps. The /cw-alert-design command reviews existing alarms for noise and coverage gaps.

Q: How does SLO management work?

A: The /cw-investigate-slo command runs a full SLO breach workflow producing an SLO Breach Explainer with burn rate, error budget remaining, impacted operations, correlated deployments, and ranked hypotheses. The /cw-slo-report command generates a portfolio-wide SLO compliance report ranked by risk.

Security & Compliance

Q: What data does Kiro Power access?

A: Kiro Power accesses only the AWS APM telemetry available through the four MCP servers: CloudWatch metrics and alarms, Application Signals service data, X-Ray traces, and CloudTrail events. It does not cache or store telemetry. CloudWatch remains the system of record.

Q: Can it modify my AWS resources?

A: By default, no. The plugin is read-only. Write actions (such as creating an alarm) require explicit in-chat confirmation and are gated by a PreToolUse hook. Destructive or billing-impacting actions (deleting log groups, modifying IAM) are never executed -- the plugin deep-links to the AWS console instead.

Q: Where is incident data stored?

A: Incident memory is stored as JSON files under .aws-apm/incidents/ on your local filesystem. No data leaves your machine or is sent to any external service. The plugin does not ship a database or backing service.

Pricing & Availability

Q: Is Kiro Power free?

A: Yes. Kiro Power is open source under the MIT license with no usage fees. Standard AWS API charges and Claude API usage from your existing subscription apply.

Q: Which AWS regions are supported?

A: Kiro Power works in any AWS region where Application Signals is available. Configure your region via the .mcp.json file or the /cw-set-context command.

INTERNAL FAQs

Technical

Q: Why build on MCP instead of calling AWS APIs directly?

A: AWS already maintains four MCP servers covering the full APM surface. Building a custom API layer would duplicate that work and create a maintenance burden. The plugin's value is the investigation method (skills, commands, hooks, artifacts, validator, memory) layered on top, not re-implementation of AWS API access.

Q: Why Cloudscape design tokens for artifact templates?

A: Cloudscape is the AWS design system. Using Cloudscape tokens ensures visual consistency with the AWS console, which is the environment SREs work in daily. The HTML artifacts render with Cloudscape styling in Cowork and degrade to structured Markdown in text-only surfaces.

Q: How do the 7 artifact templates work?

A: Each template is an HTML file with placeholder tokens (e.g., {{SLO_NAME}}, {{BURN_RATE}}). At runtime, the model fills placeholders from MCP query results. A hybrid-renderer skill provides a JSON manifest grammar for deterministic rendering. Templates include sparklines, waterfall SVGs, and interactive deep-link buttons.

Q: What is the evaluation strategy?

A: Structural tests (stdlib-only) verify manifest validity, version sync, skill/command presence, frontmatter completeness, Phase 6 presence in workflow skills, and Cloudscape token usage in templates. Investigation quality is measured via hypothesis-correctness rate against incident postmortems, time-to-first-artifact, and validator pass rate.

Business

Q: What is the target addressable market?

A: AWS has over 1 million active customers. Application Signals adoption is growing as AWS pushes its APM story. The immediate TAM is the subset of AWS customers running Application Signals with teams that use Claude Code or Cowork -- estimated at tens of thousands of SRE and DevOps practitioners in the near term.

Q: What is the competitive landscape?

A: Datadog, Honeycomb, and New Relic offer incident agents for their respective APM platforms. AWS Bedrock Agents provides a hosted agent framework. PagerDuty AIOps and Incident.io offer workflow-attached AI. Kiro Power differentiates by combining the AWS Application Signals surface with a portable plugin format and an opinionated investigation method.

Q: What is the growth strategy?

A: Phase 1: establish the plugin as the default investigation tool for Application Signals users on Claude Code. Phase 2: expand to Amazon Q Quick Apps as a remote MCP server. Phase 3: add trace-to-code developer workflows to extend from SRE investigation to developer remediation.

Q: What are the success metrics?

A: Primary: investigations per week per on-call engineer, hypothesis-correctness rate vs. postmortem root cause, time-to-first-artifact. Secondary: recurrence-flag trigger rate, validator first-pass rate, artifact share rate into incident channels.

Risks

Q: What are the main adoption barriers?

A: Application Signals must be enabled and SLOs configured before the plugin adds value (Day N tool, not Day 0). Teams not yet on Claude Code or Cowork need to adopt the agent platform first. Read-only IAM policy setup requires initial effort.

Q: How are safety concerns addressed?

A: The 5-tier action safety model ensures read-only by default. Write actions require explicit in-chat confirmation via a PreToolUse hook. Destructive and billing-impacting actions are never executed through MCP. All incident memory is local-only.

Q: What is the dependency risk on the MCP protocol?

A: MCP is an open standard with growing adoption across the AI ecosystem. The four AWS MCP servers are maintained by awslabs. If MCP evolves, the plugin's orchestration layer (skills, commands, artifacts) is independent and can adapt. Tool contracts are documented to make migration explicit.

Q: How is hallucination risk mitigated?

A: Every investigation artifact is cited to specific MCP tool calls recorded in the metadata footer. The investigation-validator runs a 6-check self-audit before any artifact is shown, verifying that every claim has a supporting evidence reference. Burn-rate and error-budget math is validated. Deep links are checked.